

Abstract:

In this updated project, the Cifar100 dataset, which consists of one hundred categories, about five hundred images per category for the training set, was classified with a convolutional neural network. With data augmentation, the test accuracy increased from 0.43 to 0.59. Transfer Learning was also implemented, which resulted in astonishing high test accuracies. In the following report, the final model's architecture and other models that have been trained are presented.

Model's Architecture:

After once again, testing several models, each with different numbers of convolutional layers and different numbers of convolutional filters and parameters, the final model's architecture consists of six convolutional layers, each followed by a max-pooling layer, five dropout layers, and two dense layers followed by an output layer.

Model: "sequential_8"

Layer (type)	Output Shape	Param #
conv2d_48 (Conv2D)	(None, 32, 32, 256)	7168
batch_normalization_64 (Batch Normalization)	(None, 32, 32, 256)	1024
conv2d_49 (Conv2D)	(None, 32, 32, 256)	590080
batch_normalization_65 (Batch Normalization)	(None, 32, 32, 256)	1024
max_pooling2d_24 (MaxPooling2D)	(None, 16, 16, 256)	0
dropout_40 (Dropout)	(None, 16, 16, 256)	0
conv2d_50 (Conv2D)	(None, 16, 16, 256)	590080
batch_normalization_66 (Batch Normalization)	(None, 16, 16, 256)	1024
conv2d_51 (Conv2D)	(None, 16, 16, 256)	590080
batch_normalization_67 (Batch Normalization)	(None, 16, 16, 256)	1024
max_pooling2d_25 (MaxPooling2D)	(None, 8, 8, 256)	0
dropout_41 (Dropout)	(None, 8, 8, 256)	0
conv2d_52 (Conv2D)	(None, 8, 8, 512)	1180160
batch_normalization_68 (Batch Normalization)	(None, 8, 8, 512)	2048
conv2d_53 (Conv2D)	(None, 8, 8, 512)	2359808
batch_normalization_69 (Batch Normalization)	(None, 8, 8, 512)	2048

max_pooling2d_26 (MaxPoolin	(None, 4, 4, 512)	0
g2D)		
dropout_42 (Dropout)	(None, 4, 4, 512)	0
flatten_8 (Flatten)	(None, 8192)	0
dense_24 (Dense)	(None, 1024)	8389632
dropout_43 (Dropout)	(None, 1024)	0
batch_normalization_70 (Bat	(None, 1024)	4096
chNormalization)		
dense_25 (Dense)	(None, 1024)	1049600
batch_normalization_71 (Bat	(None, 1024)	4096
chNormalization)		
dropout_44 (Dropout)	(None, 1024)	0
dense_26 (Dense)	(None, 100)	102500
=====		
Total params: 14,875,492		
Trainable params: 14,867,300		
Non-trainable params: 8,192		
=====		

Figure 1 The Initial Model's Architecture

The trend of changes in train and validation loss and accuracy in the range of 80 epochs is as follows.




Figure 2 Training and Validation Accuracy of Final Model




Figure 3 Training and Validation Loss of Final Model

And finally, with the use of data augmentation, the accuracy of the trained model on the test set was 59 percent.

```
Evaluate model on test data
78/78 [=====] - 2s 25ms/step - loss: 1.5819 - accuracy: 0.5913
test loss, test acc: [1.5818543434143066, 0.5913461446762085]
```

Figure 4 Accuracy of Final Model on Test Set

Transfer Learning:

In this version of the project, two pre-trained networks were used to improve the test accuracy. ResNet50 and EfficientNetB0 were imported from keras.applications.

ResNet50:


The architecture of the network used is as follows:

Model: "sequential"

Layer (type)	Output Shape	Param #
up_sampling2d (UpSampling2D)	(None, 224, 224, 3)	0
resnet50 (Functional)	(None, None, None, 2048)	23587712
global_average_pooling2d (G	(None, 2048)	0

lobalAveragePooling2D)		
dense (Dense)	(None, 512)	1049088
batch_normalization (Batch Normalization)	(None, 512)	2048
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 100)	51300
=====		
Total params: 24,690,148		
Trainable params: 1,154,532		
Non-trainable params: 23,535,616		
=====		

This model was trained for only 20 epochs due to the high number of parameters and, as a result, high training time. The trend of changes in train and validation loss and accuracy can be seen in the diagrams below.



And the test accuracy of this model is 67 percent, which can be improved if the number of epochs increases. However, because of the high number of parameters, we should be mindful of the risk of overfitting the model.

```
Evaluate model on test data
78/78 [=====] - 8s 92ms/step - loss: 1.2966 - accuracy: 0.6787
test loss, test acc: [1.2965734004974365, 0.6786859035491943]
```

Figure 5 Test Accuracy on ResNet50


EfficientNetB0:

The other pre-trained model used is EfficientNetB0, and the architecture is as follows.

Model: "sequential_2"

Layer (type)	Output Shape	Param #
up_sampling2d_2 (UpSampling2D)	(None, 224, 224, 3)	0
efficientnetb0 (Functional)	(None, None, None, 1280)	4049571
global_average_pooling2d_2 (GlobalAveragePooling2D)	(None, 1280)	0
dense_4 (Dense)	(None, 512)	655872
batch_normalization_2 (Batch Normalization)	(None, 512)	2048
dropout_2 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 100)	51300
Total params: 4,758,791		
Trainable params: 750,212		
Non-trainable params: 4,008,579		

This model was also only trained on 20 epochs. The trend of changes in train and validation loss and accuracy can be seen in the diagrams below.



The test accuracy of this model is 68 percent, which can be improved again if trained for longer.

```
Evaluate model on test data
```

```
78/78 [=====] - 6s 60ms/step - loss: 1.2337 - accuracy: 0.6847  
test loss, test acc: [1.2337055206298828, 0.6846955418586731]
```

Figure 6 Test Accuracy on EfficientNetB0